

Top 100 Interview Questions



- 1. What types of generative models have you worked with, and in what contexts?
- 2. Could you elucidate the fundamental differences between discriminative and generative models in machine learning?
- 3. Could you elaborate on the concept of generative modeling?
- 4. Could you provide examples of typical applications of generative AI?
- 5. Can you identify some well-known generative Al models?
- 6. What components are essential in a generative Al system?
- 7. How do you assess the effectiveness of a generative Al model?
- 8. Could you outline the disparities between supervised and unsupervised generative learning?
- 9. How do probability distributions contribute to generative modeling?
- 10. Can you elaborate on the concept of likelihood in generative modeling?
- 11. How do you address missing data in generative modeling processes?
- 12. What challenges do you encounter when training generative AI models?
- 13. What is mode collapse, and how does it affect generative modeling?
- 14. Could you explain adversarial training and its relevance to generative AI?
- 15. What are the distinctions between autoregressive and non-autoregressive generative models?
- 16. In what ways do attention mechanisms play a role in generative modeling?
- 17. How do you mitigate overfitting issues in generative AI?
- 18. What role do latent variables play in generative modeling?
- 19. What strategies exist for managing the diversity of generated samples?
- 20. How can generative AI be employed to create realistic images?
- 21. Can you describe the concept of variational inference in generative modeling?
- 22. What sets generative and discriminative classifiers apart?
- 23. What ethical considerations must be addressed in the development and deployment of generative AI?
- 24. How do you interpret the results produced by a generative Al model?
- 25. What limitations are associated with current generative AI techniques?
- 26. Can you explain how transfer learning is applied in generative modeling?
- 27. How do generative adversarial networks (GANs) differ from variational autoencoders (VAEs)?
- 28. What strategies do you employ when faced with imbalanced datasets in generative modeling?
- 29. What methods can be utilized to enhance the convergence of generative Al algorithms?
- 30. How do you determine the most suitable loss function for a given generative AI task?
- 31. What does "mode seeking" refer to in the context of generative modeling?

a ♦ Are you preparing for a Generative AI interview? Here are 100 possible interview questions for both junior and senior candidates:

- 32. What role do regularization techniques play in generative AI?
- 33 How do you incorporate domain knowledge into a generative Al model
- 34 Describe the difference between unconditional and conditional generative models.
- 35 Can you explain the concept of mutual information in generative modeling?
- 36 What are some techniques for generating sequences using generative AI?
- 37 How do you handle outliers in generative modeling?
- 38 Describe the difference between top-down and bottom-up generative approaches.
- 39 What are some techniques for generating natural language text using generative AI?
- 40 How do you measure the diversity of generated samples in generative modeling?
- 41 Can you explain the concept of likelihood-free inference in generative modeling?
- 42 What are some techniques for generating music using generative AI?
- 43 How do you deal with multi-modal data in generative modeling?
- 44 Describe the difference between single-modal and multi-modal generative models.
- 45 What are some techniques for generating 3D shapes using generative AI?
- 46 How do you handle sequential data in generative modeling?
- 47 Describe the difference between generative and discriminative embeddings.
- 48 What are some techniques for generating graphs using generative AI?
- 49 How do you handle continuous and discrete variables in generative modeling?
- 50 Describe the difference between generative AI and reinforcement learning.
- 51 What are some techniques for generating faces using generative AI?
- 52 How do you handle noisy data in generative modeling?
- 53 What are some techniques for generating time-series data using generative AI?
- 54 Describe the difference between generative AI and traditional machine learning.
- 55 What are some techniques for generating video using generative AI?
- 56 How do you handle variable-length sequences in generative modeling?
- 57 What are some techniques for generating medical images using generative AI?
- 58 Describe the difference between generative AI and evolutionary algorithms.
- 59 What are some techniques for generating sound using generative AI?

a ✓ Are you preparing for a Generative AI interview? Here are 100 possible interview questions for both junior and senior candidates:

- 60. Discuss the role of attention mechanisms in guiding the generation process in transformer-based generative models.
- 61. How can generative models be made interpretable by leveraging techniques like integrated gradients or concept activation vectors?
- 62. Explore the potential of probabilistic flows for efficient and diverse generation of complex data.
- 63. Discuss the concept of progressive growing and its application in training high-resolution generative models.
- 64. How can generative models be made more robust to adversarial attacks designed to fool the discriminator?
- 65. How can generative models be used for few-shot learning or domain adaptation scenarios?
- 66.Explore the potential of continual learning techniques for lifelong learning in generative models.
- 67. Discuss the role of generative models in unsupervised domain adaptation and transfer learning tasks.
- 68. How can generative models be leveraged for efficient and scalable data compression techniques?
- 69. Explain the core differences between few-shot learning and zero-shot learning approaches.
- 70. Discuss the challenges associated with training models for few-shot learning tasks.
- 71. How can metric learning techniques be leveraged to improve performance in few-shot classification tasks?
- 72. Explore the role of meta-learning approaches in adapting models to new classes with limited data (few-shot learning).
- 73. Discuss the concept of knowledge distillation and its potential benefits for fewshot learning.
- 74. How can generative models be used to augment data and improve performance in few-shot learning scenarios?
- 75. Explore the potential of transfer learning techniques for pre-training models for few-shot learning tasks.
- 76. Discuss the use of prompt-based learning for zero-shot learning tasks. How does it work?
- 77 How can we evaluate the effectiveness of zero-shot learning models?
- 78 Explain the core concept of attention mechanisms and their role in transformers architectures.
- 79. Explore the potential limitations of attention mechanisms, such as computational complexity.

- 80 Discuss recent advancements in attention mechanisms, like multi-head attention or sparse attention.
- 81 How can attention mechanisms be adapted for different tasks and data modalities (e.g., vision, text)?
- 82 Explain the core concept of Retrieval-Augmented Generation (RAG) and its benefits for question answering tasks.
- 83 Compare and contrast the functionalities of Langchain and LlamaIndex in the context of RAG systems.
- 84 How can techniques like transfer learning be applied to adapt RAG models to specific domains or question types?
- 85 How does LlamaIndex leverage vector representations of documents to enable fast and accurate document retrieval?
- 86 Discuss the different types of vector databases (e.g., Faiss, HNSW) that can be integrated with LlamaIndex for efficient document search.
- 87 How can vector databases be used beyond simple document retrieval to enhance the functionality of Generative Al models? (e.g., influencing the generation process based on retrieved vectors)
- 88 How can vector databases be scaled efficiently to handle the large data volumes typically associated with training and deploying Generative AI models?
- 89 Explore how advancements in vector database technology, such as approximate nearest neighbor search algorithms, benefit Generative AI applications.
- 90 Discuss the potential limitations of current vector database technologies and future research directions.
- 91 Discuss the limitations of supervised learning for LLM fine-tuning and how RLHF can address these limitations when human preferences are available.
- 92 How are reward signals designed for RLHF in LLM fine-tuning tasks? What factors influence the effectiveness of these reward signals?
- 93 What are some real-world applications where RLHF can be particularly beneficial for fine-tuning LLMs for specific tasks? (e.g., chatbots, dialogue systems)
- 94 Discuss the data requirements for effective LLM fine-tuning. How does the amount and quality of data impact the performance?
- 95 Explore the different strategies for freezing or fine-tuning different layers of a pre-trained LLM during the fine-tuning process.
- 96 Explain the challenges of overfitting in LLM fine-tuning and discuss different regularization techniques to mitigate it. (e.g., dropout, weight decay)
- 97 How can multi-task learning be leveraged during LLM fine-tuning to improve performance on multiple related tasks simultaneously?
- 98 What are some ethical considerations surrounding LLM fine-tuning, such as bias amplification or potential misuse of the fine-tuned model?
- 99 How can malicious actors exploit vulnerabilities in the LLM fine-tuning process to manipulate the generated outputs or introduce biases?
- 100 Discuss the potential for perpetuating or amplifying biases present in the training data during LLM fine-tuning. How can fairness metrics and techniques be integrated into LLM?

